

ATTACKING POLICY

OPEN POLICY AGENT FOR
DEVOPS ENVIRONMENTS



WWW.DEVSECOPSGUIDES.COM

Attacking Policy

• Jul 15, 2024 • 📖 22 min read

Table of contents

Allowed Repositories

- › 1. Create the ConstraintTemplate
- › 2. Create the Constraint
- › 3. Usage

Automount Service Account Token for Pod

- › 1. Create the ConstraintTemplate
- › 2. Create the Constraint
- › 3. Usage

Block Endpoint Edit Default Role

- › 1. Create the ConstraintTemplate
- › 2. Create the Constraint
- › 3. Usage

Block Services with type LoadBalancer

- › 1. Create the ConstraintTemplate
- › 2. Create the Constraint
- › 3. Usage

Block NodePort

- › ConstraintTemplate
- › Usage

Block Wildcard Ingress

- > ConstraintTemplate
- > Usage

Disallow Interactive TTY Containers

Step-by-Step Instructions

Allow Privilege Escalation in Container

Step-by-Step Instructions

Privileged Container

Read Only Root Filesystem

Host Networking Ports

App Armor

SELinux V2

Resources

Show less ^

Open Policy Agent (OPA) is a versatile tool used to enforce policies and ensure compliance within a DevSecOps environment. However, security misconfigurations in OPA can lead to significant vulnerabilities. One common issue is overly permissive policies, where misconfigured rules allow more access than intended. This can result in unauthorized access to sensitive data or critical services. Additionally, inadequate logging and monitoring of policy decisions can obscure the detection of policy violations and security incidents, making it difficult for security teams to respond promptly. Ensuring that policies are correctly scoped, thoroughly tested, and continuously monitored is crucial to maintaining a secure DevSecOps environment.

Another critical aspect of OPA security misconfiguration is the improper integration with other DevSecOps tools and services. For instance, if OPA is not correctly

integrated with Kubernetes or CI/CD pipelines, it can lead to gaps in policy enforcement, allowing insecure code or configurations to be deployed. Moreover, the lack of proper version control and review processes for policy changes can introduce vulnerabilities, as unreviewed changes might weaken security postures. Therefore, it is essential to establish a robust governance framework for managing policy changes, coupled with automated testing and validation to ensure that security policies remain effective and aligned with organizational security standards. Regular audits and updates to policies are necessary to adapt to evolving threats and ensure comprehensive security coverage.

Allowed Repositories

To enforce policies that restrict container images to specific repositories using Open Policy Agent (OPA) and Gatekeeper, follow these steps. The following example will demonstrate how to create a ConstraintTemplate and Constraint to ensure container images begin with a specified list of prefixes.

1. Create the ConstraintTemplate

First, define the ConstraintTemplate. This template describes the policy logic and the parameters it requires.

Create a file named `k8sallowedrepos_template.yaml` with the following content:

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sallowedrepos
  annotations:
    metadata.gatekeeper.sh/title: "Allowed Repositories"
    metadata.gatekeeper.sh/version: 1.0.1
  description: >-
    Requires container images to begin with a string from the
    specified list.
```

```
spec:  
  crd:  
    spec:  
      names:  
        kind: K8sAllowedRepos  
      validation:  
        # Schema for the `parameters` field  
      openAPIV3Schema:  
        type: object  
        properties:  
          repos:  
            description: The list of prefixes a container image is  
allowed to have.  
            type: array  
            items:  
              type: string  
targets:  
  - target: admission.k8s.gatekeeper.sh  
    rego: |  
      package k8sallowedrepos  
  
      violation[{"msg": msg}] {  
        container := input.review.object.spec.containers[_]  
        not startswith_any(container.image, input.parameters.repos)  
        msg := sprintf("container <%v> has an invalid image repo  
<%v>, allowed repos are %v", [container.name, container.image,  
input.parameters.repos])  
      }  
  
      violation[{"msg": msg}] {  
        container := input.review.object.spec.initContainers[_]  
        not startswith_any(container.image, input.parameters.repos)  
        msg := sprintf("initContainer <%v> has an invalid image repo  
<%v>, allowed repos are %v", [container.name, container.image,  
input.parameters.repos])  
      }
```

```
violation[{"msg": msg}] {
    container := input.review.object.spec.ephemeralContainers[_]
    not startswith_any(container.image, input.parameters.repos)
    msg := sprintf("ephemeralContainer <%v> has an invalid image
repo <%v>, allowed repos are %v", [container.name, container.image,
input.parameters.repos])
}

startswith_any(image, repos) {
    some repo
    startswith(image, repo)
}
```

Apply the template to your Kubernetes cluster:

COPY 

```
kubectl apply -f k8sallowedrepos_template.yaml
```

2. Create the Constraint

Next, define a Constraint that uses the ConstraintTemplate to specify the allowed repository prefixes.

Create a file named `k8sallowedrepos_constraint.yaml` with the following content:

COPY 

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sAllowedRepos
metadata:
  name: allowed-repos
spec:
  match:
    kinds:
      - apiGroups: []
        kinds: ["Pod"]
```

```
parameters:  
  repos:  
    - "myregistry.io/myrepo"  
    - "anotherregistry.io/anotherrepo"
```

Apply the constraint to your Kubernetes cluster:

```
kubectl apply -f k8sallowedrepos_constraint.yaml
```

COPY 

3. Usage

To test this setup, try deploying a Pod with an image that does not match the allowed repository prefixes:

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: invalid-image-pod  
spec:  
  containers:  
    - name: myapp  
      image: invalidregistry.io/myapp:latest
```

COPY 

Save this as `invalid-image-pod.yaml` and attempt to apply it:

```
kubectl apply -f invalid-image-pod.yaml
```

COPY 

Automount Service Account Token for Pod

To control the ability of any Pod to enable `automountServiceAccountToken` using Open Policy Agent (OPA) and Gatekeeper, you can create a ConstraintTemplate and a Constraint. This setup will enforce policies that prevent Pods from automatically mounting service account tokens unless explicitly allowed.

1. Create the ConstraintTemplate

First, define the ConstraintTemplate. This template describes the policy logic and the parameters it requires.

Create a file named `k8spspautomountserviceaccounttokenpod_template.yaml` with the following content:

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8spspautomountserviceaccounttokenpod
  annotations:
    metadata.gatekeeper.sh/title: "Automount Service Account Token for Pod"
    metadata.gatekeeper.sh/version: 1.0.1
    description: >-
      Controls the ability of any Pod to enable automountServiceAccountToken.
spec:
  crd:
    spec:
      names:
        kind: K8sPSPAutomountServiceAccountTokenPod
      validation:
        openAPIV3Schema:
          type: object
          description: >-
            Controls the ability of any Pod to enable automountServiceAccountToken.
      targets:
```

```
- target: admission.k8s.gatekeeper.sh  
  rego: |  
    package
```

Apply the template to your Kubernetes cluster:

COPY 

```
kubectl apply -f k8spspautomountserviceaccounttokenpod_template.yaml
```

2. Create the Constraint

Next, define a Constraint that uses the ConstraintTemplate to specify the enforcement policy.

Create a file named `k8spspautomountserviceaccounttokenpod_constraint.yaml` with the following content:

COPY 

```
apiVersion: constraints.gatekeeper.sh/v1beta1  
kind: K8sPSPAutomountServiceAccountTokenPod  
metadata:  
  name: disallow-automount-sa-token  
spec:  
  match:  
    kinds:  
      - apiGroups: [""]  
        kinds: ["Pod"]
```

Apply the constraint to your Kubernetes cluster:

COPY 

```
kubectl apply -
```

```
f k8spsautomountserviceaccounttokenpod_constraint.yaml=
```

3. Usage

To test this setup, try deploying a Pod with `automountServiceAccountToken` enabled:

COPY 

```
apiVersion: v1
kind: Pod
metadata:
  name: automount-sa-token-pod
spec:
  automountServiceAccountToken: true
  containers:
    - name: myapp
      image: nginx:latest
```

Save this as `automount-sa-token-pod.yaml` and attempt to apply it:

COPY 

```
kubectl apply -f automount-sa-token-pod.yaml
```

Block Endpoint Edit Default Role

To prevent the `system:aggregate-to-edit` ClusterRole from granting permissions to create, patch, or update Endpoints in Kubernetes, you can create a ConstraintTemplate and a Constraint using Open Policy Agent (OPA) and Gatekeeper. This setup addresses the security concern outlined in CVE-2021-25740.

1. Create the ConstraintTemplate

First, define the ConstraintTemplate. This template describes the policy logic and the parameters it requires.

Create a file named `k8sblockendpointeditdefaultrole_template.yaml` with the following content:

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sblockendpointeditdefaultrole
  annotations:
    metadata.gatekeeper.sh/title: "Block Endpoint Edit Default Role"
    metadata.gatekeeper.sh/version: 1.0.0
    description: >-
      Many Kubernetes installations by default have a
      system:aggregate-to-edit
      ClusterRole which does not properly restrict access to editing
      Endpoints.

      This ConstraintTemplate forbids the system:aggregate-to-edit
      ClusterRole
      from granting permission to create/patch/update Endpoints.

      ClusterRole/system:aggregate-to-edit should not allow
      Endpoint edit permissions due to CVE-2021-25740, Endpoint &
      EndpointSlice
      permissions allow cross-Namespace forwarding,
      https://github.com/kubernetes/kubernetes/issues/103675

spec:
  crd:
    spec:
      names:
        kind: K8sBlockEndpointEditDefaultRole
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8sblockendpointeditdefaultrole

        violation[{"msg": msg}] {
          input.review.object.metadata.name == "system:aggregate-to-
```

```
edit"
    endpointRule(input.review.object.rules[_])
        msg := "ClusterRole system:aggregate-to-edit should not
allow endpoint edit permissions. For k8s version < 1.22, the Cluster
Role should be annotated with
rbac.authorization.kubernetes.io/autoupdate=false to prevent
autoreconciliation back to default permissions for this role."
}

endpointRule(rule) {
    "endpoints" == rule.resources[_]
    hasEditVerb(rule.verbs)
}

hasEditVerb(verbs) {
    "create" == verbs[_]
}

hasEditVerb(verbs) {
    "patch" == verbs[_]
}

hasEditVerb(verbs) {
    "update" == verbs[_]
}
```

Apply the template to your Kubernetes cluster:

COPY 

```
kubectl apply -f k8sblockendpointeditdefaultrole_template.yaml
```

2. Create the Constraint

Next, define a Constraint that uses the ConstraintTemplate to specify the enforcement policy.

Create a file named `k8sblockendpointeditdefaultrole_constraint.yaml` with the following content:

COPY 

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sBlockEndpointEditDefaultRole
metadata:
  name: block-endpoint-edit-default-role
spec:
  match:
    kinds:
      - apiGroups: [""]
        kinds: ["ClusterRole"]
```

Apply the constraint to your Kubernetes cluster:

COPY 

```
kubectl apply -f k8sblockendpointeditdefaultrole_constraint.yaml
```

3. Usage

To test this setup, try applying a ClusterRole named `system:aggregate-to-edit` with permissions to edit Endpoints:

COPY 

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: system:aggregate-to-edit
rules:
  - apiGroups: [""]
    resources: ["endpoints"]
    verbs: ["create", "patch", "update"]
```

Save this as `test-clusterrole.yaml` and attempt to apply it:

COPY 

```
kubectl apply -f test-clusterrole.yaml
```

Block Services with type LoadBalancer

To disallow all Services with type `LoadBalancer` in your Kubernetes cluster using Open Policy Agent (OPA) and Gatekeeper, you can create a `ConstraintTemplate` and a `Constraint`. This setup will enforce policies that prevent the creation of `LoadBalancer`-type Services, adhering to the best practices for network security.

1. Create the ConstraintTemplate

First, define the `ConstraintTemplate`. This template describes the policy logic.

Create a file named `k8sblockloadbalancer_template.yaml` with the following content:

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sblockloadbalancer
  annotations:
    metadata.gatekeeper.sh/title: "Block Services with type
LoadBalancer"
    metadata.gatekeeper.sh/version: 1.0.0
  description: >-
    Disallows all Services with type LoadBalancer.
```

<https://kubernetes.io/docs/concepts/services-networking/service/#loadbalancer>

```
spec:
  crd:
    spec:
```

```
names:
  kind: K8sBlockLoadBalancer
targets:
  - target: admission.k8s.gatekeeper.sh
    rego: |
      package k8sblockloadbalancer

      violation[{"msg": msg}] {
        input.review.kind.kind == "Service"
        input.review.object.spec.type == "LoadBalancer"
        msg := "User is not allowed to create service of type
LoadBalancer"
      }
```

Apply the template to your Kubernetes cluster:

COPY 

```
kubectl apply -f k8sblockloadbalancer_template.yaml
```

2. Create the Constraint

Next, define a Constraint that uses the ConstraintTemplate to specify the enforcement policy.

Create a file named `k8sblockloadbalancer_constraint.yaml` with the following content:

COPY 

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sBlockLoadBalancer
metadata:
  name: block-loadbalancer-services
spec:
  match:
    kinds:
```

```
- apiGroups: []
  kinds: ["Service"]
```

Apply the constraint to your Kubernetes cluster:

COPY 

```
kubectl apply -f k8sblockloadbalancer_constraint.yaml
```

3. Usage

To test this setup, try creating a Service with type `LoadBalancer`:

COPY 

```
apiVersion: v1
kind: Service
metadata:
  name: test-loadbalancer-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

Save this as `test-loadbalancer-service.yaml` and attempt to apply it:

COPY 

```
kubectl apply -f test-loadbalancer-service.yaml
```

The OPA Gatekeeper should deny the creation of the Service and return a violation message specifying that creating a Service of type LoadBalancer is not allowed.

Block NodePort

To block the creation of Services with type `NodePort` in your Kubernetes cluster using Open Policy Agent (OPA) and Gatekeeper, you can use a ConstraintTemplate. This setup will enforce a policy that disallows the creation of `NodePort` type Services.

ConstraintTemplate

Define the ConstraintTemplate. This template describes the policy logic.

Create a file named `k8sblocknodeport_template.yaml` with the following content:

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sblocknodeport
  annotations:
    metadata.gatekeeper.sh/title: "Block NodePort"
    metadata.gatekeeper.sh/version: 1.0.0
  description: >-
    Disallows all Services with type NodePort.

    https://kubernetes.io/docs/concepts/services-networking/service/#nodeport
spec:
  crd:
    spec:
      names:
        kind: K8sBlockNodePort
  targets:
  - target: admission.k8s.gatekeeper.sh
    rego: |
```

```
package k8sblocknodeport

violation[{"msg": msg}] {
    input.review.kind.kind == "Service"
    input.review.object.spec.type == "NodePort"
    msg := "User is not allowed to create service of type
NodePort"
}
```

Apply the template to your Kubernetes cluster:

COPY 

```
kubectl apply -f k8sblocknodeport_template.yaml
```

Usage

You can apply this ConstraintTemplate directly from the Gatekeeper library:

COPY 

```
kubectl apply -f https://raw.githubusercontent.com/open-
policy-agent/gatekeeper-library/master/library/general/block-
nodeport-services/template.yaml
```

Block Wildcard Ingress

To block the creation of Ingresses with a blank or wildcard (*) hostname in your Kubernetes cluster using Open Policy Agent (OPA) and Gatekeeper, you can use a ConstraintTemplate. This setup ensures that users cannot create Ingresses that can intercept traffic for other services in the cluster.

ConstraintTemplate

Define the ConstraintTemplate. This template describes the policy logic.

Create a file named `k8sblockwildcardingress_template.yaml` with the following content:

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sblockwildcardingress
  annotations:
    metadata.gatekeeper.sh/title: "Block Wildcard Ingress"
    metadata.gatekeeper.sh/version: 1.0.1
    description: >-
      Users should not be able to create Ingresses with a blank or
      wildcard (*) hostname since that would enable them to intercept
      traffic for other services in the cluster, even if they don't have
      access to those services.
spec:
  crd:
    spec:
      names:
        kind: K8sBlockWildcardIngress
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package K8sBlockWildcardIngress

        contains_wildcard(hostname) = true {
          hostname == ""
        }

        contains_wildcard(hostname) = true {
          contains(hostname, "*")
        }

      violation[{"msg": msg}] {
        input.review.kind.kind == "Ingress"
        # object.get is required to detect omitted host fields
        hostname := object.get(input.review.object.spec.rules[_],
```

```
"host", "")  
    contains_wildcard(hostname)  
    msg := sprintf("Hostname '%v' is not allowed since it counts  
as a wildcard, which can be used to intercept traffic from other  
applications.", [hostname])  
}
```

Apply the template to your Kubernetes cluster:

```
kubectl apply -f k8sblockwildcardingress_template.yaml
```

Usage

You can apply this ConstraintTemplate directly from the Gatekeeper library:

```
kubectl apply -f https://raw.githubusercontent.com/open-policy-agent/gatekeeper-library/master/library/general/block-wildcard-ingress/template.yaml
```

Disallow Interactive TTY Containers

when applying the `Disallow Interactive TTY Containers` ConstraintTemplate using Gatekeeper, you can follow these instructions:

Step-by-Step Instructions

1. Create ConstraintTemplate YAML File:

Save the following YAML content into a file named

```
k8sdisallowinteractivetty_template.yaml.
```

COPY 

```
apiVersion: templates.gatekeeper.sh/v1  
kind: ConstraintTemplate  
metadata:  
  name: k8sdisallowinteractivetty  
  annotations:
```

```
metadata.gatekeeper.sh/title: "Disallow Interactive TTY  
Containers"  
metadata.gatekeeper.sh/version: 1.0.0  
description: >--  
    Requires that objects have the fields `spec.tty` and  
`spec.stdin` set to false or unset.  
spec:  
  crd:  
    spec:  
      names:  
        kind: K8sDisallowInteractiveTTY  
      validation:  
        openAPIV3Schema:  
          type: object  
          properties:  
            exemptImages:  
              description: >--  
                Any container that uses an image that matches an entry  
in this list will be excluded  
                from enforcement. Prefix-matching can be signified with  
`*`. For example: `my-image-*`.  
  
                It is recommended that users use the fully-qualified  
Docker image name (e.g. start with a domain name)  
                in order to avoid unexpectedly exempting images from an  
untrusted repository.  
              type: array  
              items:  
                type: string  
targets:  
- target: admission.k8s.gatekeeper.sh  
  rego: |  
    package k8sdisallowinteractivetty  
  
    import data.lib.exempt_container.is_exempt  
  
    violation[{"msg": msg, "details": {}}] {
```

```
c := input_containers[_]
not is_exempt(c)
input_allow_interactive_fields(c)
msg := sprintf("Containers using tty or stdin (%v) are not
allowed running image: %v", [c.name, c.image])
}

input_allow_interactive_fields(c) {
    has_field(c, "stdin")
    not c.stdin == false
}
input_allow_interactive_fields(c) {
    has_field(c, "tty")
    not c.tty == false
}
input_containers[c] {
    c := input.review.object.spec.containers[_]
}
input_containers[c] {
    c := input.review.object.spec.ephemeralContainers[_]
}
input_containers[c] {
    c := input.review.object.spec.initContainers[_]
}
has_field(object, field) = true {
    object[field]
}
libs:
- |
    package lib.exempt_container

is_exempt(container) {
    exempt_images := object.get(object.get(input,
"parameters", {}), "exemptImages", [])
    img := container.image
    exemption := exempt_images[_]
    _matches_exemption(img, exemption)
```

```
        }
    }

    _matches_exemption(img, exemption) {
        not endswith(exemption, "*")
        exemption == img
    }

    _matches_exemption(img, exemption) {
        endswith(exemption, "*")
        prefix := trim_suffix(exemption, "*")
        startswith(img, prefix)
    }
}
```

Apply the ConstraintTemplate:

Use `kubectl apply` to apply the ConstraintTemplate directly from its URL:

COPY 

```
kubectl apply -f https://raw.githubusercontent.com/open-
policy-agent/gatekeeper-
library/master/library/general/disallowinteractive/template.yaml
```

This command will deploy the ConstraintTemplate `k8sdisallowinteractive` into your Kubernetes cluster.

Allow Privilege Escalation in Container

when applying the `Allow Privilege Escalation in Container` ConstraintTemplate using Gatekeeper, follow these steps:

Step-by-Step Instructions

1. Create ConstraintTemplate YAML File:

Save the following YAML content into a file named

k8spspallowprivilegeescalationcontainer_template.yaml.

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8spspallowprivilegeescalationcontainer
  annotations:
    metadata.gatekeeper.sh/title: "Allow Privilege Escalation in Container"
    metadata.gatekeeper.sh/version: 1.1.0
  description: >-
    Controls restricting escalation to root privileges. Corresponds to the
    `allowPrivilegeEscalation` field in a PodSecurityPolicy. For more
    information, see
    https://kubernetes.io/docs/concepts/policy/pod-security-policy/#privilege-escalation
spec:
  crd:
    spec:
      names:
        kind: K8sPSPAllowPrivilegeEscalationContainer
      validation:
        openAPIV3Schema:
          type: object
          description: >-
            Controls restricting escalation to root privileges.
Corresponds to the
  `allowPrivilegeEscalation` field in a PodSecurityPolicy. For more
  information, see
  https://kubernetes.io/docs/concepts/policy/pod-security-policy/#privilege-escalation
  properties:
```

```
exemptImages:
  description: >-
    Any container that uses an image that matches an entry
    in this list will be excluded
      from enforcement. Prefix-matching can be signified with
      `*`. For example: `my-image-*`.

      It is recommended that users use the fully-qualified
    Docker image name (e.g. start with a domain name)
      in order to avoid unexpectedly exempting images from an
    untrusted repository.

  type: array
  items:
    type: string
targets:
  - target: admission.k8s.gatekeeper.sh
    code:
      - engine: K8sNativeValidation
        source:
          variables:
            - name: containers
              expression: 'has(variables.anyObject.spec.containers) ? variables.anyObject.spec.containers : []'
            - name: initContainers
              expression:
                'has(variables.anyObject.spec.initContainers) ? variables.anyObject.spec.initContainers : []'
            - name: ephemeralContainers
              expression:
                'has(variables.anyObject.spec.ephemeralContainers) ? variables.anyObject.spec.ephemeralContainers : []'
            - name: exemptImagePrefixes
              expression: |
                !has(variables.params.exemptImages) ? [] :
                  variables.params.exemptImages.filter(image,
                    image.endsWith("*")).map(image, string(image).replace("*", ""))
            - name: exemptImageExplicit
```

```
expression: |
    !has(variables.params.exemptImages) ? [] :
        variables.params.exemptImages.filter(image,
!image.endsWith("*"))
    - name: exemptImages
        expression: |
            (variables.containers + variables.initContainers +
variables.ephemeralContainers).filter(container,
                container.image in variables.exemptImageExplicit
||

                variables.exemptImagePrefixes.exists(exemption,
string(container.image).startsWith(exemption))
                    ).map(container, container.image)
    - name: badContainers
        expression: |
            (variables.containers + variables.initContainers +
variables.ephemeralContainers).filter(container,
                !(container.image in variables.exemptImages) && (
                    !has(container.securityContext) ||


!has(container.securityContext.allowPrivilegeEscalation) ||
container.securityContext.allowPrivilegeEscalation != false
                )
            )
    validations:
        - expression: '(has(request.operation) &&
request.operation == "UPDATE") || size(variables.badContainers) == 0'
            messageExpression: '"Privilege escalation container is
not allowed: " + variables.badContainers.map(c, c.image).join(", ")'
        - engine: Rego
            source:
                rego: |
                    package k8spspallowprivilegescalationcontainer

import data.lib.exclude_update.is_update
import data.lib.exempt_container.is_exempt
```

```
violation[{"msg": msg, "details": {}}] {
    #
spec.containers.securityContext.allowPrivilegeEscalation field is
immutable.
        not is_update(input.review)

        c := input_containers[_]
        not is_exempt(c)
        input_allow_privilege_escalation(c)
        msg := sprintf("Privilege escalation container is
not allowed: %v", [c.name])
    }

input_allow_privilege_escalation(c) {
    not has_field(c, "securityContext")
}
input_allow_privilege_escalation(c) {
    not c.securityContext.allowPrivilegeEscalation ==
false
}
input_containers[c] {
    c := input.review.object.spec.containers[_]
}
input_containers[c] {
    c := input.review.object.spec.initContainers[_]
}
input_containers[c] {
    c := input.review.object.spec.ephemeralContainers[_]
}
# has_field returns whether an object has a field
has_field(object, field) = true {
    object[field]
}
libs:
- |
    package lib.exclude_update
```

```
is_update(review) {
    review.operation == "UPDATE"
}
-
|
package lib.exempt_container

is_exempt(container) {
    exempt_images := object.get(object.get(input,
"parameters", {}), "exemptImages", [])
    img := container.image
    exemption := exempt_images[_]
    _matches_exemption(img, exemption)
}

_matches_exemption(img, exemption) {
    not endswith(exemption, "*")
    exemption == img
}

_matches_exemption(img, exemption) {
    endswith(exemption, "*")
    prefix := trim_suffix(exemption, "*")
    startswith(img, prefix)
}
```

Apply the ConstraintTemplate:

Use `kubectl apply` to apply the ConstraintTemplate directly from its URL:

COPY 

```
kubectl apply -f https://raw.githubusercontent.com/open-
policy-agent/gatekeeper-library/master/library/pod-
security-policy/allow-privilege-escalation/template.yaml
```

Privileged Container

when applying the `Privileged Container` ConstraintTemplate using Gatekeeper, follow these steps:

1. Create ConstraintTemplate YAML File:

Save the following YAML content into a file named

`k8spspprivilagedcontainer_template.yaml`.

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8spspprivilagedcontainer
  annotations:
    metadata.gatekeeper.sh/title: "Privileged Container"
    metadata.gatekeeper.sh/version: 1.1.0
    description: >-
      Controls the ability of any container to enable privileged mode.
      Corresponds to the `privileged` field in a PodSecurityPolicy.
```

For more

information, see

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/#privileged>

spec:

crd:

spec:

names:

kind: K8sPSPPrivilagedContainer

validation:

openAPIV3Schema:

type: object

description: >-

Controls the ability of any container to enable privileged mode.

Corresponds to the `privileged` field in a

PodSecurityPolicy. For more

```
information, see
https://kubernetes.io/docs/concepts/policy/pod-security-
policy/#privileged

properties:

  exemptImages:
    description: >-
      Any container that uses an image that matches an entry
      in this list will be excluded
      from enforcement. Prefix-matching can be signified with
      `*`. For example: `my-image-*`.

      It is recommended that users use the fully-qualified
      Docker image name (e.g. start with a domain name)
      in order to avoid unexpectedly exempting images from an
      untrusted repository.

    type: array
    items:
      type: string

targets:
  - target: admission.k8s.gatekeeper.sh

  code:
    - engine: K8sNativeValidation

    source:
      variables:
        - name: containers
          expression: 'has(variables.anyObject.spec.containers)'

? variables.anyObject.spec.containers : []
  - name: initContainers
    expression:
      'has(variables.anyObject.spec.initContainers) ?'
      variables.anyObject.spec.initContainers : []
        - name: ephemeralContainers
          expression:
            'has(variables.anyObject.spec.ephemeralContainers) ?'
            variables.anyObject.spec.ephemeralContainers : []
              - name: exemptImagePrefixes
                expression: |
```

```
!has(variables.params.exemptImages) ? [] :  
    variables.params.exemptImages.filter(image,  
image.endsWith("*")).map(image, string(image).replace("*", ""))  
    - name: exemptImageExplicit  
        expression: |  
            !has(variables.params.exemptImages) ? [] :  
                variables.params.exemptImages.filter(image,  
!image.endsWith("*"))  
                - name: exemptImages  
                    expression: |  
                        (variables.containers + variables.initContainers +  
variables.ephemeralContainers).filter(container,  
            container.image in variables.exemptImageExplicit  
||  
            variables.exemptImagePrefixes.exists(exemption,  
string(container.image).startsWith(exemption)))  
            - name: badContainers  
                expression: |  
                    (variables.containers + variables.initContainers +  
variables.ephemeralContainers).filter(container,  
            !(container.image in variables.exemptImages) &&  
            (has(container.securityContext) &&  
has(container.securityContext.privileged) &&  
container.securityContext.privileged == true)  
            ).map(container, "Privileged container is not  
allowed: " + container.name +", securityContext: " +  
container.securityContext)  
                validations:  
                - expression: '(has(request.operation) &&  
request.operation == "UPDATE") || size(variables.badContainers) == 0'  
                    messageExpression:  
'variables.badContainers.join("\n")'  
- engine: Rego  
source:  
rego: |  
    package k8spspprivilaged
```

```
import data.lib.exclude_update.is_update
import data.lib.exempt_container.is_exempt

violation[{"msg": msg, "details": {}}] {
    # spec.containers.privileged field is immutable.
    not is_update(input.review)

    c := input_containers[_]
    not is_exempt(c)
    c.securityContext.privileged
    msg := sprintf("Privileged container is not allowed:
%v, securityContext: %v", [c.name, c.securityContext])
}

input_containers[c] {
    c := input.review.object.spec.containers[_]
}

input_containers[c] {
    c := input.review.object.spec.initContainers[_]
}

input_containers[c] {
    c := input.review.object.spec.ephemeralContainers[_]
}

libs:
- |
    package lib.exclude_update

        is_update(review) {
            review.operation == "UPDATE"
        }
- |
    package lib.exempt_container

        is_exempt(container) {
            exempt_images := object.get(object.get(input,
```

```
"parameters", {}), "exemptImages", [])
    img := container.image
    exemption := exempt_images[_]
    _matches_exemption(img, exemption)
}

_matches_exemption(img, exemption) {
    not endswith(exemption, "*")
    exemption == img
}

_matches_exemption(img, exemption) {
    endswith(exemption, "*")
    prefix := trim_suffix(exemption, "*")
    startswith(img, prefix)
}
```

Apply the ConstraintTemplate:

Use `kubectl apply` to apply the ConstraintTemplate directly from its URL:

COPY 

```
kubectl apply -f https://raw.githubusercontent.com/open-
policy-agent/gatekeeper-library/master/library/pod-
security-policy/privileged-containers/template.yaml
```

Read Only Root Filesystem

the `Read Only Root Filesystem` ConstraintTemplate using Gatekeeper and skip steps 2 and 3, follow these instructions:

1. Create ConstraintTemplate YAML File:

Save the following YAML content into a file named

```
k8spspreadonlyrootfilesystem_template.yaml.
```

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8spspreadonlyrootfilesystem
  annotations:
    metadata.gatekeeper.sh/title: "Read Only Root Filesystem"
    metadata.gatekeeper.sh/version: 1.1.0
  description: >-
    Requires the use of a read-only root file system by pod
    containers.
    Corresponds to the `readOnlyRootFilesystem` field in a
    PodSecurityPolicy. For more information, see
    https://kubernetes.io/docs/concepts/policy/pod-security-policy/#volumes-and-file-systems
spec:
  crd:
    spec:
      names:
        kind: K8sPSPReadOnlyRootFilesystem
      validation:
        openAPIV3Schema:
          type: object
          description: >-
            Requires the use of a read-only root file system by pod
            containers.
            Corresponds to the `readOnlyRootFilesystem` field in a
            PodSecurityPolicy. For more information, see
            https://kubernetes.io/docs/concepts/policy/pod-security-policy/#volumes-and-file-systems
          properties:
            exemptImages:
              description: >-
                Any container that uses an image that matches an entry
                in this list will be excluded
                from enforcement. Prefix-matching can be signified with
                `*`. For example: `my-image-*`.
```

It is recommended that users use the fully-qualified Docker image name (e.g. start with a domain name) in order to avoid unexpectedly exempting images from an untrusted repository.

```
    type: array
    items:
        type: string
targets:
- target: admission.k8s.gatekeeper.sh
code:
- engine: K8sNativeValidation
source:
variables:
- name: containers
    expression: 'has(variables.anyObject.spec.containers) ? variables.anyObject.spec.containers : []'
- name: initContainers
    expression:
'has(variables.anyObject.spec.initContainers) ? variables.anyObject.spec.initContainers : []'
- name: ephemeralContainers
    expression:
'has(variables.anyObject.spec.ephemeralContainers) ? variables.anyObject.spec.ephemeralContainers : []'
- name: exemptImagePrefixes
    expression: |
        !has(variables.params.exemptImages) ? [] :
            variables.params.exemptImages.filter(image,
image.endsWith("*")).map(image, string(image).replace("*", ""))
- name: exemptImageExplicit
    expression: |
        !has(variables.params.exemptImages) ? [] :
            variables.params.exemptImages.filter(image,
!image.endsWith("*"))
- name: exemptImages
    expression: |
```

```
(variables.containers + variables.initContainers +
variables.ephemeralContainers).filter(container,
    container.image in variables.exemptImageExplicit
||

        variables.exemptImagePrefixes.exists(exemption,
string(container.image).startsWith(exemption)))
    - name: badContainers
        expression: |
            (variables.containers + variables.initContainers +
variables.ephemeralContainers).filter(container,
                !(container.image in variables.exemptImages) &&
                (!has(container.securityContext) ||


!has(container.securityContext.readOnlyRootFilesystem) ||
                container.securityContext.readOnlyRootFilesystem
!= true)
                    ).map(container, container.name)
validations:
    - expression: '(has(request.operation) &&
request.operation == "UPDATE") || size(variables.badContainers) == 0'
        messageExpression: '"only read-only root filesystem
container is allowed: " + variables.badContainers.join(", ")'

    - engine: Rego
        source:
            rego: |
                package k8spspreadonlyrootfilesystem

                import data.lib.exclude_update.is_update
                import data.lib.exempt_container.is_exempt

                violation[{"msg": msg, "details": {}}] {
                    # spec.containers.readOnlyRootFilesystem field is
immutable.

                    not is_update(input.review)

                    c := input_containers[_]
```

```
    not is_exempt(c)
    input_read_only_root_fs(c)
    msg := sprintf("only read-only root filesystem
container is allowed: %v", [c.name])
}

input_read_only_root_fs(c) {
    not has_field(c, "securityContext")
}
input_read_only_root_fs(c) {
    not c.securityContext.readOnlyRootFilesystem == true
}

input_containers[c] {
    c := input.review.object.spec.containers[_]
}
input_containers[c] {
    c := input.review.object.spec.initContainers[_]
}
input_containers[c] {
    c := input.review.object.spec.ephemeralContainers[_]
}

# has_field returns whether an object has a field
has_field(object, field) = true {
    object[field]
}
libs:
- |
    package lib.exclude_update

    is_update(review) {
        review.operation == "UPDATE"
    }
- |
    package lib.exempt_container
```

```
is_exempt(container) {
    exempt_images := object.get(object.get(input,
"parameters", {}), "exemptImages", [])
    img := container.image
    exemption := exempt_images[_]
    _matches_exemption(img, exemption)
}

_matches_exemption(img, exemption) {
    not endswith(exemption, "*")
    exemption == img
}

_matches_exemption(img, exemption) {
    endswith(exemption, "*")
    prefix := trim_suffix(exemption, "*")
    startswith(img, prefix)
}
```

Apply the ConstraintTemplate:

Use `kubectl apply` to apply the ConstraintTemplate directly from its URL:

COPY 

```
kubectl apply -f https://raw.githubusercontent.com/open-
policy-agent/gatekeeper-library/master/library/pod-
security-policy/read-only-root-filesystem/template.yaml
```

Host Networking Ports

the `Host Networking Ports` ConstraintTemplate using Gatekeeper and skip steps 2 and 3, follow these instructions:

1. Create ConstraintTemplate YAML File:

Save the following YAML content into a file named

k8spsphostnetworkingports_template.yaml.

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8spsphostnetworkingports
  annotations:
    metadata.gatekeeper.sh/title: "Host Networking Ports"
    metadata.gatekeeper.sh/version: 1.1.1
    description: >-
      Controls usage of host network namespace by pod containers.
      HostNetwork verification happens without exception for exemptImages.
      Specific
      ports must be specified. Corresponds to the `hostNetwork` and
      `hostPorts` fields in a PodSecurityPolicy. For more information,
      see
      https://kubernetes.io/docs/concepts/policy/pod-security-policy/#host-namespaces
spec:
  crd:
    spec:
      names:
        kind: K8sPSPHostNetworkingPorts
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          type: object
          description: >-
            Controls usage of host network namespace by pod containers.
            HostNetwork verification happens without exception for exemptImages.
            Specific
            ports must be specified. Corresponds to the `hostNetwork` and
            `hostPorts` fields in a PodSecurityPolicy. For more
            information, see
```

```
https://kubernetes.io/docs/concepts/policy/pod-security-
policy/#host-namespaces

properties:
  exemptImages:
    description: >-
      Any container that uses an image that matches an entry
      in this list will be excluded
      from enforcement. Prefix-matching can be signified with
      `*`. For example: `my-image-*`.
```

It is recommended that users use the fully-qualified Docker image name (e.g. start with a domain name)

in order to avoid unexpectedly exempting images from an untrusted repository.

```
type: array
  items:
    type: string
  hostNetwork:
    description: "Determines if the policy allows the use of HostNetwork in the pod spec."
    type: boolean
    min:
      description: "The start of the allowed port range,
      inclusive."
      type: integer
    max:
      description: "The end of the allowed port range,
      inclusive."
      type: integer
  targets:
    - target: admission.k8s.gatekeeper.sh
      code:
        - engine: K8sNativeValidation
          source:
            variables:
              - name: containers
                expression: 'has(variables.anyObject.spec.containers)'
```

```
? variables.anyObject.spec.containers : []
  - name: initContainers
    expression:
'has(variables.anyObject.spec.initContainers) ?'
variables.anyObject.spec.initContainers : []
  - name: ephemeralContainers
    expression:
'has(variables.anyObject.spec.ephemeralContainers) ?'
variables.anyObject.spec.ephemeralContainers : []
  - name: exemptImagePrefixes
    expression: |
      !has(variables.params.exemptImages) ? [] :
        variables.params.exemptImages.filter(image,
image.endsWith("*")).map(image, string(image).replace("*", ""))
  - name: exemptImageExplicit
    expression: |
      !has(variables.params.exemptImages) ? [] :
        variables.params.exemptImages.filter(image,
!image.endsWith("*"))
  - name: exemptImages
    expression: |
      (variables.containers + variables.initContainers +
variables.ephemeralContainers).filter(container,
        container.image in variables.exemptImageExplicit
||
      variables.exemptImagePrefixes.exists(exemption,
string(container.image).startsWith(exemption)))
  - name: badContainers
    expression: |
      (variables.containers + variables.initContainers +
variables.ephemeralContainers).filter(container,
        !(container.image in variables.exemptImages) &&
        (
          (container.ports.all(port, has(port.hostPort) &&
has(variables.params.min) && port.hostPort < variables.params.min)) ||
          (container.ports.all(port, has(port.hostPort) &&
has(variables.params.max) && port.hostPort > variables.params.max))
```

```
        )
    )

validations:
  - expression: '(has(request.operation) &&
request.operation == "UPDATE") || size(variables.badContainers) == 0'
    messageExpression: '"The specified hostNetwork and
hostPort are not allowed, pod: " + variables.anyObject.metadata.name +
". Allowed values: " + variables.params'
  - expression: |
    (has(request.operation) && request.operation ==
"UPDATE") ||
    (!has(variables.params.hostNetwork) ||
!variables.params.hostNetwork ?
(has(variables.anyObject.spec.hostNetwork) &&
!variables.anyObject.spec.hostNetwork) : true)
    messageExpression: '"The specified hostNetwork and
hostPort are not allowed, pod: " + variables.anyObject.metadata.name +
". Allowed values: " + variables.params'
  - engine: Rego
  source:
    rego: |
      package k8spsphostnetworkingports

      import data.lib.exclude_update.is_update
      import data.lib.exempt_container.is_exempt

      violation[{"msg": msg, "details": {}}] {
        # spec.hostNetwork field is immutable.
        not is_update(input.review)

        input_share_hostnetwork(input.review.object)
        msg := sprintf("The specified hostNetwork and
hostPort are not allowed, pod: %v. Allowed values: %v",
[input.review.object.metadata.name, input.parameters])
      }

      input_share_hostnetwork(o) {
```

```
        not input.parameters.hostNetwork
        o.spec.hostNetwork
    }

    input_share_hostnetwork(_) {
        hostPort := input_containers[_].ports[_].hostPort
        hostPort < input.parameters.min
    }

    input_share_hostnetwork(_) {
        hostPort := input_containers[_].ports[_].hostPort
        hostPort > input.parameters.max
    }

    input_containers[c] {
        c := input.review.object.spec.containers[_]
        not is_exempt(c)
    }

    input_containers[c] {
        c := input.review.object.spec.initContainers[_]
        not is_exempt(c)
    }

    input_containers[c] {
        c :=
input.review.object.spec.ephemeralContainers[_]
        not is_exempt(c)
    }
libs:
- |
    package lib.exclude_update

    is_update(review) {
        review.operation == "UPDATE"
    }
- |
```

```
package lib.exempt_container

    is_exempt(container) {
        exempt_images := object.get(object.get(input,
"parameters", {}), "exemptImages", [])
        img := container.image
        exemption := exempt_images[_]
        _matches_exemption(img, exemption)
    }

    _matches_exemption(img, exemption) {
        not endswith(exemption, "*")
        exemption == img
    }

    _matches_exemption(img, exemption) {
        e
```

Apply the ConstraintTemplate:

Use `kubectl apply` to apply the ConstraintTemplate directly from its URL:

COPY 

```
kubectl apply -f https://raw.githubusercontent.com/open-
policy-agent/gatekeeper-library/master/library/pod-
security-policy/host-network-ports/template.yaml
```

This command will deploy the ConstraintTemplate `k8spsphostnetworkingports` into your Kubernetes cluster.

App Armor

the `App Armor` ConstraintTemplate using Gatekeeper, follow these streamlined instructions:

1. Create ConstraintTemplate YAML File:

Save the following YAML content into a file named `k8spapparmor_template.yaml`.

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8spapparmor
  annotations:
    metadata.gatekeeper.sh/title: "App Armor"
    metadata.gatekeeper.sh/version: 1.0.0
    description: >-
      Configures an allow-list of AppArmor profiles for use by
      containers.

      This corresponds to specific annotations applied to a
      PodSecurityPolicy.

      For information on AppArmor, see
      https://kubernetes.io/docs/tutorials/clusters/apparmor/

spec:
  crd:
    spec:
      names:
        kind: K8sPSPAppArmor
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          type: object
          description: >-
            Configures an allow-list of AppArmor profiles for use by
            containers.

            This corresponds to specific annotations applied to a
            PodSecurityPolicy.

            For information on AppArmor, see
            https://kubernetes.io/docs/tutorials/clusters/apparmor/
          properties:
            exemptImages:
              description: >-
```

Any container that uses an image that matches an entry in this list will be excluded from enforcement. Prefix-matching can be signified with `*`. For example: `my-image-*`.

It is recommended that users use the fully-qualified Docker image name (e.g. start with a domain name)

in order to avoid unexpectedly exempting images from an untrusted repository.

```
type: array
items:
  type: string
allowedProfiles:
  description: "An array of AppArmor profiles. Examples: `runtime/default`, `unconfined`."
  type: array
  items:
    type: string
targets:
- target: admission.k8s.gatekeeper.sh
rego: |
  package k8spspapparmor

  import data.lib.exempt_container.is_exempt

  violation[{"msg": msg, "details": {}}] {
    metadata := input.review.object.metadata
    container := input_containers[_]
    not is_exempt(container)
    not input_apparmor_allowed(container, metadata)
    msg := sprintf("AppArmor profile is not allowed, pod: %v,
container: %v. Allowed profiles: %v",
[input.review.object.metadata.name, container.name,
input.parameters.allowedProfiles])
  }

  input_apparmor_allowed(container, metadata) {
```

```
    get_annotation_for(container, metadata) ==
input.parameters.allowedProfiles[_]
}

input_containers[c] {
    c := input.review.object.spec.containers[_]
}
input_containers[c] {
    c := input.review.object.spec.initContainers[_]
}
input_containers[c] {
    c := input.review.object.spec.ephemeralContainers[_]
}

get_annotation_for(container, metadata) = out {
    out =
metadata.annotations[sprintf("container.apparmor.security.beta.kubernetes.io/%v",
[container.name])]
}
get_annotation_for(container, metadata) = out {
    not
metadata.annotations[sprintf("container.apparmor.security.beta.kubernetes.io/%v",
[container.name])]
    out = "runtime/default"
}
libs:
- |
    package lib.exempt_container

    is_exempt(container) {
        exempt_images := object.get(object.get(input,
"parameters", {}), "exemptImages", [])
        img := container.image
        exemption := exempt_images[_]
        _matches_exemption(img, exemption)
    }
}
```

```
_matches_exemption(img, exemption) {
    not endswith(exemption, "*")
    exemption == img
}

_matches_exemption(img, exemption) {
    endswith(exemption, "*")
    prefix := trim_suffix(exemption, "*")
    startswith(img, prefix)
}
```

Apply the ConstraintTemplate:

Use `kubectl apply` to apply the ConstraintTemplate directly from its URL

COPY 

```
kubectl apply -f https://raw.githubusercontent.com/open-
policy-agent/gatekeeper-library/master/library/pod-
security-policy/apparmor/template.yaml
```

SELinux V2

the `SELinux V2` ConstraintTemplate using Gatekeeper while skipping steps 2 and 3, follow these streamlined instructions:

1. Create ConstraintTemplate YAML File:

Save the following YAML content into a file named

`k8spspselinuxv2_template.yaml`.

COPY 

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8spspselinuxv2
```

```
annotations:
  metadata.gatekeeper.sh/title: "SELinux V2"
  metadata.gatekeeper.sh/version: 1.0.3
  description: >-
    Defines an allow-list of seLinuxOptions configurations for pod
    containers. Corresponds to a PodSecurityPolicy requiring SELinux
    configs.

    For more information, see
      https://kubernetes.io/docs/concepts/policy/pod-security-policy/#selinux

spec:
  crd:
    spec:
      names:
        kind: K8sPSPSELinuxV2
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          type: object
          description: >-
            Defines an allow-list of seLinuxOptions configurations for
            pod
            containers. Corresponds to a PodSecurityPolicy requiring
            SELinux configs.

            For more information, see
              https://kubernetes.io/docs/concepts/policy/pod-security-policy/#selinux

            properties:
              exemptImages:
                description: >-
                  Any container that uses an image that matches an entry
                  in this list will be excluded
                  from enforcement. Prefix-matching can be signified with
                  `*`. For example: `my-image-*`.

                  It is recommended that users use the fully-qualified
                  Docker image name (e.g. start with a domain name)
```

in order to avoid unexpectedly exempting images from an untrusted repository.

```
  type: array
  items:
    type: string
  allowedSELinuxOptions:
    type: array
    description: "An allow-list of SELinux options
configurations."
    items:
      type: object
      description: "An allowed configuration of SELinux
options for a pod container."
      properties:
        level:
          type: string
          description: "An SELinux level."
        role:
          type: string
          description: "An SELinux role."
        type:
          type: string
          description: "An SELinux type."
        user:
          type: string
          description: "An SELinux user."
targets:
- target: admission.k8s.gatekeeper.sh
  rego: |
    package k8spspselinux

    import data.lib.exclude_update.is_update
    import data.lib.exempt_container.is_exempt

    # Disallow top level custom SELinux options
    violation[{"msg": msg, "details": {}}] {
      # spec.securityContext.selinuxOptions field is immutable.
```

```
    not is_update(input.review)

        has_field(input.review.object.spec.securityContext,
"seLinuxOptions")
            not
input_seLinuxOptions_allowed(input.review.object.spec.securityContext.
seLinuxOptions)
            msg := sprintf("SELinux options is not allowed, pod: %v.
Allowed options: %v", [input.review.object.metadata.name,
input.parameters.allowedSELinuxOptions])
        }
# Disallow container level custom SELinux options
violation[{"msg": msg, "details": {}}] {
    # spec.containers.securityContext.seLinuxOptions field is
immutable.
        not is_update(input.review)

        c := input_security_context[_]
        not is_exempt(c)
        has_field(c.securityContext, "seLinuxOptions")
        not
input_seLinuxOptions_allowed(c.securityContext.seLinuxOptions)
            msg := sprintf("SELinux options is not allowed, pod: %v,
container: %v. Allowed options: %v",
[input.review.object.metadata.name, c.name,
input.parameters.allowedSELinuxOptions])
        }
input_seLinuxOptions_allowed(options) {
    params := input.parameters.allowedSELinuxOptions[_]
    field_allowed("level", options, params)
    field_allowed("role", options, params)
    field_allowed("type", options, params)
    field_allowed("user", options, params)
}
field_allowed(field, options, params) {
```

```
params[field] == options[field]
}

field_allowed(field, options, _) {
    not has_field(options, field)
}

input_security_context[c] {
    c := input.review.object.spec.containers[_]
    has_field(c.securityContext, "seLinuxOptions")
}

input_security_context[c] {
    c := input.review.object.spec.initContainers[_]
    has_field(c.securityContext, "seLinuxOptions")
}

input_security_context[c] {
    c := input.review.object.spec.ephemeralContainers[_]
    has_field(c.securityContext, "seLinuxOptions")
}

# has_field returns whether an object has a field
has_field(object, field) = true {
    object[field]
}

libs:
- |
    package lib.exclude_update

        is_update(review) {
            review.operation == "UPDATE"
        }

- |
    package lib.exempt_container

        is_exempt(container) {
            exempt_images := object.get(object.get(input,
"parameters", {}), "exemptImages", [])
            img := container.image
        }
}
```

```
exemption := exempt_images[_]
_matches_exemption(img, exemption)
}

_matches_exemption(img, exemption) {
    not endswith(exemption, "*")
    exemption == img
}

_matches_exemption(img, exemption) {
    endswith(exemption, "*")
    prefix := trim_suffix(exemption, "*")
    startswith(img, prefix)
}
```

Apply the ConstraintTemplate:

Use `kubectl apply` to apply the ConstraintTemplate directly from its URL:

COPY 

```
kubectl apply -f https://raw.githubusercontent.com/open-
policy-agent/gatekeeper-library/master/library/pod-
security-policy/selinux/template.yaml
```

Resources

- <https://open-policy-agent.github.io/gatekeeper-library/website/validation/selinux/>
- <https://devsecopsguides.com>
- <https://www.openpolicyagent.org/integrations>

Written by



Reza Rashidi



Add your bio

Published on



DevSecOpsGuides



Add blog description

MORE ARTICLES



Reza Rashidi



Attacking IaC

Attacking Infrastructure as Code (IaC) methods involves exploiting vulnerabilities and misconfigurat...



Reza Rashidi



Attacking Vagrant

Vagrant, a tool for building and managing virtual machine environments, is widely used for developme...



Reza Rashidi



Attacking Golang

Golang (or Go) is a statically typed, compiled programming language designed at Google. It is known ...

©2024 DevSecOpsGuides

[Archive](#) · [Privacy policy](#) · [Terms](#)



Write on Hashnode

Powered by [Hashnode](#) - Home for tech writers and readers